# CSC321H5 Project 3.

Deadline: Thursday, March. 19, by 9pm

Submission: Submit a PDF export of the completed notebook.

Late Submission: Please see the syllabus for the late submission criteria.

In this assignment, we will build a convolutional neural network that can predict whether two shoes are from the **same pair** or from two **different pairs**. This kind of application can have real-world applications: for example to help people who are visually impaired to have more independence.

We will explore two convolutional architectures. While we will give you starter code to help make data processing a bit easier, you'll have a chance to build your neural network all by yourself!

You may modify the starter code as you see fit, including changing the signatures of functions and adding/removing helper functions. However, please make sure that your TA can understand what you are doing and why.

If you find exporting the Google Colab notebook to be difficult, you can create your own PDF report that includes your code, written solutions, and outputs that the graders need to assess your work.

```
In [26]:
```

```
import pandas
import numpy as np
import matplotlib.pyplot as plt

import time
import random

import torch
import torch.nn as nn
import torch.optim as optim
```

## Question 1. Data

Download the data from the course website at <a href="https://www.cs.toronto.edu/~lczhang/321/files/p3data.zip">https://www.cs.toronto.edu/~lczhang/321/files/p3data.zip</a>

Unzip the file. There are three main folders: train, test\_w and test\_m. Data in train will be used for training and validation, and the data in the other folders will be used for testing. This is so that the entire class will have the same test sets.

We've separated test\_w and test\_m so that we can track our model performance for women's shoes and men's shoes separately. Each of the test sets contain images from 10 students who submitted images of either exclusively men's shoes or women's shoes.

Upload this data to Google Colab. Then, mount Google Drive from your Google Colab notebook:

```
In [ ]:
```

```
from google.colab import drive
drive.mount('/content/gdrive')
```

After you have done so, read this entire section (ideally this entire handout) before proceeding. There are right and wrong ways of processing this data. If you don't make the correct choices, you may find yourself needing to start over. Many machine learning projects fail because of the lack of care taken during the data processing stage.

# Part (a) -- 1 pts

Why might we care about the accuracies of the men's and women's shoes as two separate measures? Why would we expect our model accuracies for the two groups to be different?

Recall that your application may help people who are visually impaired.

```
In [53]:
```

```
# Your answer goes here. Please make sure it is not cut off
"""

By keeping track of a different test accuracy for men's and women's shoes,
we are able to determine if our model has overfit to identify men's or women's
shoes due to the amount of men's or women's shoes present in our training data
By having one be higher than the other, we can correct an overfit by evening
out the ratio of men's to women's shoes in the training data or tune the hyperparameters
to account for this difference.
"""
```

### Out[53]:

"\nBy keeping track of a different test accuracy for men's and women's shoes,\nwe are able to dete rmine if our model is overfit to identify men's or women's\nshoes due to the amount of men's or wo men's shoes present in our training data\nwhen we were training the model. By having one be higher than the other, we \ncan correct an overfit by evening out the ratio of men's to women's shoes in\nthe training data.\n"

# Part (b) -- 4 pts

Load the training and test data, and separate your training data into training and validation. Create the numpy arrays  $train_data$ ,  $valid_data$ ,  $test_w$  and  $test_m$ , all of which should be of shape [\*, 3, 2, 224, 224, 3]. The dimensions of these numpy arrays are as follows:

- \* the number of students allocated to train, valid, or test
- 3 the 3 pairs of shoes submitted by that student
- 2 the left/right shoes
- 224 the height of each image
- 224 the width of each image
- 3 the colour channels

So, the item  $train_{data[4,0,0,:,::]}$  should give us the left shoe of the first image submitted by the 5th student. The item  $train_{data[4,0,1,:,::]}$  should be the right shoe in the same pair. The item  $train_{data[4,1,1,:,::]}$  should be the right shoe in a different pair, submitted by the same student.

When you first load the images using (for example) plt.imread, you may see a numpy array of shape [224, 224, 4] instead of [224, 224, 3]. That last channel is the alpha channel for transparent pixels, and should be removed. The pixel intensities are stored as an integer between 0 and 255. Divide the intensities by 255 so that you have floating-point values between 0 and 1. Then, subtract 0.5 so that the elements of train\_data, valid\_data and test\_data are between -0.5 and 0.5. Note that this step actually makes a huge difference in training!

This function might take a while to run---it takes 3-4 minutes for me to just load the files from Google Drive. If you want to avoid running this code multiple times, you can save your numpy arrays and load it later: <a href="https://docs.scipy.org/doc/numpy/reference/generated/numpy.save.html">https://docs.scipy.org/doc/numpy/reference/generated/numpy.save.html</a>

### In [40]:

```
# Your code goes here. Make sure it does not get cut off
# You can use the code below to help you get started. You're welcome to modify
# the code or remove it entirely: it's just here so that you don't get stuck
# reading files
import glob
path = "C:/Users/aaron/Documents/School/CSC321/Projects/Project 3/data/train/*.jpg" # edit me
images = []
# we have a total of 678 images in the training folder, 678/2 = 339 pairs of shoes, 339/3 = 113 st
udents
# we will use the last 11 students' images out of our 113 total for the validation set (last 66/67
8)
for file in glob.glob(path):
   filename = file.split("/")[-1]
                                    # get the name of the .jpg file
   img = plt.imread(file)
                                    # read the image as a numpy array
   img = img[:, :, :3] # remove the alpha channel
   img = (img/255)-0.5 # apply modifications from instructions
   images.append(img)
train data = []
valid data = []
index = 0
for student in range(int((len(images))/6)):
```

```
submission = []
    for pair in range(3):
        shoes = []
        for side in range(2):
            shoes.append(images[index])
            index+=1
        submission.append(np.stack(shoes, axis=0))
    if index >= 611:
       valid data.append(np.stack(submission, axis=0))
    else:
       train data.append(np.stack(submission, axis=0))
valid data = np.array(valid data)
train_data = np.array(train_data)
# for test m
path = "C:/Users/aaron/Documents/School/CSC321/Projects/Project 3/data/test m/*.jpg" # edit me
images = []
for file in glob.glob(path):
   filename = file.split("/")[-1] # get the name of the .jpg file
    img = plt.imread(file)
                                     # read the image as a numpy array
   img = img[:, :, :3] # remove the alpha channel
   img = (img/255)-0.5 \# apply modifications from instructions
   images.append(img)
test m = []
index = 0
for student in range(int((len(images))/6)):
    submission = []
   for pair in range(3):
       shoes = []
        for side in range(2):
           shoes.append(images[index])
            index+=1
        submission.append(np.stack(shoes, axis=0))
   test_m.append(np.stack(submission, axis=0))
test m = np.array(test m)
# for test w
path = "C:/Users/aaron/Documents/School/CSC321/Projects/Project 3/data/test w/*.jpg" # edit me
images = []
for file in glob.glob(path):
   filename = file.split("/")[-1] # get the name of the .jpg file
   img = plt.imread(file)
                                    # read the image as a numpy array
   img = img[:, :, :3] # remove the alpha channel
   img = (img/255)-0.5 \# apply modifications from instructions
   images.append(img)
test w = []
index = 0
for student in range(int((len(images))/6)):
    submission = []
   for pair in range(3):
       shoes = []
        for side in range(2):
            shoes.append(images[index])
            index+=1
        submission.append(np.stack(shoes, axis=0))
   test w.append(np.stack(submission, axis=0))
test w = np.array(test w)
```

## In [37]:

```
# Run this code, include the image in your PDF submission
plt.figure()
plt.imshow(train_data[4,0,0,:,:,:]) # left shoe of first pair submitted by 5th student
plt.figure()
plt.imshow(train_data[4,0,1,:,:,:]) # right shoe of first pair submitted by 5th student
plt.figure()
plt.imshow(train_data[4,1,1,:,:,:]) # right shoe of second pair submitted by 5th student

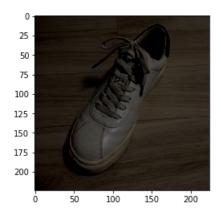
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
```

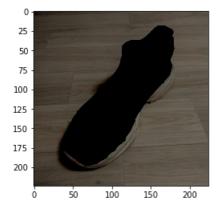
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

Out[37]:

<matplotlib.image.AxesImage at 0x278831ad198>







# Part (c) -- 2 pts

Since we want to train a model that determines whether two shoes come from the **same** pair or **different** pairs, we need to create some labelled training data. Our model will take in an image, either consisting of two shoes from the **same pair** or from **different pairs**. So, we'll need to generate some *positive examples* with images containing two shoes that *are* from the same pair, and some *negative examples* where images containing two shoes that *are not* from the same pair. We'll generate the *positive examples* in this part, and the *negative examples* in part (c).

Write a function <code>generate\_same\_pair()</code> that takes one of the data sets that you produced in part (a), and generates a numpy array where each pair of shoes in the data set is concatenated together. In particular, we'll be concatenating together images of left and right shoes along the <code>height</code> axis. Your function <code>generate\_same\_pair</code> should return a numpy array of shape <code>[\*, 448, 224, 3]</code>.

(Later on, we will need to convert this numpy array into a PyTorch tensor with shape [\*, 3, 448, 224] . For now, we'll keep the RGB channel as the last dimension since that's what plt.imshow requires)

### In [6]:

```
# Your code goes here
def generate_same_pair(data):
    out = np.ndarray((np.shape(data)[0]*3, 448, 224, 3))
    for student in range(np.shape(data)[0]):
        for shoes in range(3):
            out[student*3 + shoes, :224, :, :] = data[student, shoes, 0, :, :, :]
            out[student*3 + shoes, 224:, :, :] = data[student, shoes, 1, :, :, :]
    return out

# Run this code, include the result with your PDF submission
print(train_data.shape) # if this is [N, 3, 2, 224, 224, 3]
print(generate_same_pair(train_data).shape) # should be [N*3, 448, 224, 3]
plt.imshow(generate_same_pair(train_data)[0]) # should show 2 shoes from the same pair
```

```
(101, 3, 2, 224, 224, 3)
(303, 448, 224, 3)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

### Out[6]:

<matplotlib.image.AxesImage at 0x27880307390>



# Part (d) -- 2 pts

Write a function <code>generate\_different\_pair()</code> that takes one of the data sets that you produced in part (a), and generates a numpy array in the same shape as part (b). However, each image will contain 2 shoes from a **different** pair, but submitted by the **same student**. Do this by jumbling the 3 pairs of shoes submitted by each student.

Theoretically, for each student image submissions, there are 6 different combinations of "wrong pairs" that we could produce. To keep our data set *balanced*, we will only produce **three** combinations of wrong pairs per unique person. In other words, generate same pairs and generate different pairs should return the same number of training examples.

### In [7]:

```
# Your code goes here
def generate_different_pair(data):
    out = np.ndarray((np.shape(data)[0]*3, 448, 224, 3))
    for student in range(np.shape(data)[0]):
        for shoes in range(3):
            out[student*3 + shoes, :224, :, :] = data[student, shoes, 0, :, :, :]
            out[student*3 + shoes, 224:, :, :] = data[student, (shoes+1)%3, 1, :, :, :]
    return out

# Run this code, include the result with your PDF submission
print(train_data.shape) # if this is [N, 3, 2, 224, 224, 3]
print(generate_different_pair(train_data).shape) # should be [N*3, 448, 224, 3]
plt.imshow(generate_different_pair(train_data)[0]) # should show 2 shoes from different pairs
```

```
(101, 3, 2, 224, 224, 3)
(303, 448, 224, 3)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

#### Out[7]:

<matplotlib.image.AxesImage at 0x2788036aef0>



# Part (e) -- 1 pts

Why do we insist that the different pairs of shoes still come from the same student? (Hint: what else do images from the same student have in common?)

#### In [54]:

```
# Your answer goes here. Please make sure it is not cut off
"""

We insist that the different pairs of shoes still come from the same student because this way we a
re able to
(hopefully) keep the same background features so that the biggest difference between each image is
the shoe and
not anything in the background like the lighting or texture. Also, each student cannot upload the
same pair of
shoes twice, so we can gurantee that we can generate different pairs from their data.
"""
```

## Out[54]:

'\nWe insist that the different pairs of shoes still come from the same student because this way we are able to\n(hopefully) keep the same background features so that the biggest difference between each image is the shoe and\nnot anything in the background like the lighting or texture.\n'

# Part (f) -- 1 pts

Why is it important that our data set be *balanced*? In other words suppose we created a data set where 99% of the images are of shoes that are *not* from the same pair, and 1% of the images are shoes that *are* from the same pair. Why could this be a problem?

# In [56]:

```
# Your answer goes here. Please make sure it is not cut off
"""

It is important that our data set be balanced because if we had a data set where 99% of
the images are of shoes that are not from the same pair, and 1% of the images are shoes
that are from the same pair, our model would be more inclined to predict that two given
shoes are from a different pair rather than the same pair.
"""
```

## Out[56]:

'\nIt is important that our data set be balanced because if we had a data set where 99% of\nthe im ages are of shoes that are not from the same pair, and 1% of the images are shoes\nthat are from the same pair, our model would be more inclined to predict that two given\nshoes are from a different pair rather than the same pair.\n'

# **Question 2. Convolutional Neural Networks**

Before starting this question, we recommend reviewing the lecture and tutorial materials on convolutional neural networks.

In this section, we will build two CNN models in PyTorch.

## Part (a) -- 4 pts

Implement a CNN model in PyTorch called CNN that will take images of size  $3 \times 448 \times 224$ , and classify whether the images contain shoes from the same pair or from different pairs.

The model should contain the following layers:

- A convolution layer that takes in 3 channels, and outputs n channels.
- A  $2 \times 2$  downsampling (either using a strided convolution in the previous step, or max pooling)
- A second convolution layer that takes in n channels, and outputs  $n \times 2$  channels.
- ullet A 2 imes2 downsampling (either using a strided convolution in the previous step, or max pooling)
- A third convolution layer that takes in n imes 2 channels, and outputs n imes 4 channels.
- A  $2 \times 2$  downsampling (either using a strided convolution in the previous step, or max pooling)
- A fourth convolution layer that takes in  $n \times 4$  channels, and outputs  $n \times 8$  channels.
- ullet A 2 imes 2 downsampling (either using a strided convolution in the previous step, or max pooling)
- · A fully-connected layer with 100 hidden units
- · A fully-connected layer with 2 hidden units

Make the variable n a parameter of your CNN. You can use either  $3 \times 3$  or  $5 \times 5$  convolutions kernels. Set your padding to be (kernel size - 1) / 2 so that your feature maps have an even height/width.

Note that we are omitting certain steps that practitioners will typically not mention, like ReLU activations and reshaping operations. Use the tutorial materials and your past projects to figure out where they are.

### In [10]:

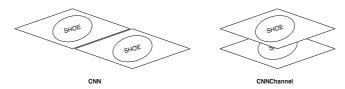
```
class CNN (nn.Module):
        __init__(self, n=4):
   def
       super(CNN, self).__init__()
        self.n = n
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=n, kernel_size=5, stride=2, padding=2)
        self.conv2 = nn.Conv2d(in channels=n, out channels=n*2, kernel size=5, stride=2, padding=2)
        self.conv3 = nn.Conv2d(in channels=n*2, out channels=n*4, kernel size=5, stride=2, padding=2
       self.conv4 = nn.Conv2d(in channels=n*4, out channels=n*8, kernel size=5, stride=2, padding=2
        self.fc1 = nn.Linear(in features=14*28*8*n, out features=100)
        self.fc2 = nn.Linear(in features=100, out features=2)
   def forward(self, inp):
       x = inp.transpose(2, 3)
       c1 = torch.relu(self.conv1(x))
       c2 = torch.relu(self.conv2(c1))
       c3 = torch.relu(self.conv3(c2))
       c4 = torch.relu(self.conv4(c3))
       c5 = torch.relu(self.fc1(c4.flatten(1, 3)))
       c6 = self.fc2(c5)
        return c6
```

# Part (b) -- 4 pts

Implement a CNN model in PyTorch called CNNChannel that contains the same layers as in the Part (a), but with one crucial difference: instead of starting with an image of shape  $3\times448$ , we will first manipulate the image so that the left and right shoes

< 224

images are concatenated along the channel dimension.



Complete the manipulation in the <code>forward()</code> method (by slicing and using the function <code>torch.cat</code>). The input to the first convolutional layer should have 6 channels instead of 3 (input shape  $6 \times 224$ ).

 $\times 224$ 

Use the same hyperparameter choices as you did in part (a), e.g. for the kernel size, choice of downsampling, and other choices.

#### In [14]:

```
class CNNChannel (nn.Module) :
   def init (self, n=4):
       super(CNNChannel, self). init ()
       self.n = n
        self.conv1 = nn.Conv2d(in_channels=6, out_channels=n, kernel_size=5, stride=2, padding=2)
        self.conv2 = nn.Conv2d(in channels=n, out channels=n*2, kernel size=5, stride=2, padding=2)
        self.conv3 = nn.Conv2d(in channels=n*2, out channels=n*4, kernel size=5, stride=2, padding=2
        self.conv4 = nn.Conv2d(in channels=n*4, out channels=n*8, kernel size=5, stride=2, padding=2
        self.fc1 = nn.Linear(in features=14*14*8*n, out features=100)
        self.fc2 = nn.Linear(in features=100, out features=2)
   def forward(self, inp):
       x = inp.transpose(2, 3)
       left = x[:, :, :224, :]
       right = x[:, :, 224:, :]
       merged = torch.cat((left, right), dim=1)
       c1 = torch.relu(self.conv1(merged))
       c2 = torch.relu(self.conv2(c1))
       c3 = torch.relu(self.conv3(c2))
       c4 = torch.relu(self.conv4(c3))
       c5 = torch.relu(self.fc1(c4.flatten(1, 3)))
       c6 = self.fc2(c5)
       return c6
```

# Part (c) -- 2 pts

Although our task is a binary classification problem, we will still use the architecture of a multi-class classification problem. That is, we'll use a one-hot vector to represent our target (just like in Project 2). We'll also use CrossEntropyLoss instead of BCEWithLogitsLoss. In fact, this is a standard practice in machine learning because this architecture performs better!

Explain why this architecture might give you better performance.

```
In [57]:
```

```
# Your answer goes here. Please make sure it is not cut off
"""

This architecture might give a better performance because with our one-hot vector already in place
to represent our target, we can save on computing time by simply using cross entropy loss.

BCEWithLogitsLoss requires us to apply a sigmoid function to our output before calculating
cross entropy loss, so it uses more computing power. This also means that it is prone to numerical
instability.
"""
```

### Out[57]:

'\nThis architecture might give a better performance because with our one-hot vector already in pl ace\nto represent our target, we can save on computing time by simply using cross entropy loss.\nBCEWithLogitsLoss requires us to apply a sigmoid function to our output before calculating\ncross entropy loss, so it uses more computing power. This also means that it is prone to numerical\ninstability.\n'

# Part (d) -- 2 pts

The two models are quite similar, and should have almost the same number of parameters. However, one of these models will perform better, showing that architecture choices **do** matter in machine learning. Explain why one of these models performs better.

```
In [58]:
```

```
# Your answer goes here. Please make sure it is not cut off
"""

The CNNChannel model would likely perform better because the extra parameters in the
```

```
model will allow us to more finely tune the inputs on average while the we are training the model to get a better validation accuracy.
```

### Out[58]:

'\nThe CNNChannel model would likely perform better because the extra parameters in the\nmodel will allow us to more finely tune the inputs on average while the we are training\nthe model to ge t a better validation accuracy.\n'

# Part (e) -- 2 pts

The function <code>get\_accuracy</code> is written for you. You may need to modify this function depending on how you set up your model and training.

Unlike in project 2, we will separately compute the model accuracy on the positive and negative samples. Explain why we may wish to track these two values separately.

### In [59]:

```
# Your answer goes here. Please make sure it is not cut off
"""

We may wish to track these two values separately because the positive
model accuracy will represent the model correctly identifying shoes
from the same pair while the negative value will represent the model
whenever it identifies shoes of different pairs. It will be important
for the model to be able to predict both cases, so it will need a set
of data that it can compare new data to in order to determine if a
newly introduced pair of shoes are of the same or different pairs.
"""
```

### Out [59]:

'\nWe may wish to track these two values separately because the positive\nmodel accuracy will represent the model correctly identifying shoes\nfrom the same pair while the negative value will represent the model\nwhenever it identifies shoes of different pairs. It will be important\nfor the model to be able to predict both cases so it will need a set\nof data that it can compare new data to in order to determine if a\nnewly introduced pair of shoes are of the same or different pairs.\n'

## In [19]:

```
def get accuracy(model, data, batch size=50):
    """Compute the model accuracy on the data set. This function returns two
   separate values: the model accuracy on the positive samples,
   and the model accuracy on the negative samples.
   Example Usage:
   >>> model = CNN() # create untrained model
   >>> pos_acc, neg_acc= get_accuracy(model, valid_data)
   >>> false_positive = 1 - pos_acc
   >>> false negative = 1 - neg acc
   model.eval()
   n = data.shape[0]
   data pos = generate same pair(data)
                                        # should have shape [n * 3, 448, 224, 3]
   data_neg = generate_different_pair(data) # should have shape [n * 3, 448, 224, 3]
   pos correct = 0
   for i in range(0, len(data pos), batch size):
       xs = torch.Tensor(data_pos[i:i+batch_size]).transpose(1, 3)
       zs = model(xs)
       pred = zs.max(1, keepdim=True)[1] # get the index of the max logit
       pred = pred.detach().numpy()
       pos correct += (pred == 1).sum()
   neg correct = 0
   for i in range(0, len(data_neg), batch_size):
    xs = torch.Tensor(data neg[i:i+batch size]).transpose(1, 3)
```

```
zs = model(xs)
    pred = zs.max(1, keepdim=True)[1] # get the index of the max logit
   pred = pred.detach().numpy()
   neg correct += (pred == 0).sum()
return pos_correct / (n * 3), neg_correct / (n * 3)
```

# **Question 3. Training**

Now, we will write the functions required to train the model.

## Part (a) -- 10 pts

Write the function train model that takes in (as parameters) the model, training data, validation data, and other hyperparameters like the batch size, weight decay, etc. This function should be somewhat similar to the training code that you wrote in Project 2, but with a major difference in the way we treat our training data.

Since our positive and negative training sets are separate, it is actually easier for us to generate separate minibatches of positive and negative training data! In each iteration, we'll take batch size / 2 positive samples and batch size / 2 negative samples. We will also generate labels of 1's for the positive samples, and 0's for the negative samples.

Here's what we will be looking for:

- · main training loop; choice of loss function; choice of optimizer
- obtaining the positive and negative samples
- shuffling the positive and negative samples at the start of each epoch
- in each iteration, take batch size / 2 positive samples and batch size / 2 negative samples as our input for this
- in each iteration, take np.ones(batch size / 2) as the labels for the positive samples, and np.zeros(batch size / 2) as the labels for the negative samples
- · conversion from numpy arrays to PyTorch tensors, making sure that the input has dimensions "NCHW", use the .transpose() method in either PyTorch or numpy
- · computing the forward and backward passes
- after every epoch, checkpoint your model (Project 2 has in-depth instructions and examples for how to do this)
- after every epoch, report the accuracies for the training set and validation set
- · track the training curve information and plot the training curve

## In [20]:

```
# Write your code here
def make data(data):
   same = generate same pair(data)
   dif = generate_different_pair(data)
   out1 = []
   out2 = []
   for i in range (np.shape (same) [0]):
       out1.append(same[i, :, :, :])
   for i in range(np.shape(dif)[0]):
       out2.append(dif[i, :, :, :])
   return out1, out2
def get batch(same, dif, start, end):
   xs = []
   for i in range(start//2, end//2):
       xs.append(same[i])
   for i in range(start//2, end//2):
       xs.append(dif[i])
   return xs, np.concatenate((np.ones((end-start)//2), np.zeros((end-start)//2)))
def train model (model, t data, v data, batch size=50, learning rate=0.001, num epochs=100, checkpoi
nt path='./ckpt-epoch{}.pk'):
   criterion = nn.CrossEntropyLoss()
   optimizer = optim.Adam (model.parameters(), lr=learning rate)
   optimizer.zero grad()
   iters, losses, train_acc, val_acc, iters_sub, t_pos, t_negs, v_pos, v_negs = [], [], [], []
, [], [], [], []
   same, dif = make data(t data)
```

```
n = 0
    start = time.time()
    last = start
    print("")
    for epoch in range(num epochs):
        random.shuffle(same)
        random.shuffle(dif)
        for batch in range(0, np.shape(same)[0], batch size):
            if (batch + batch_size) > np.shape(same)[0]:
            xs, ts = get batch(same, dif, batch, batch + batch size)
            xs = torch.Tensor(xs)
            xs = xs.transpose(1, 3)
            ts = torch.Tensor(ts).long()
            zs = model(xs)
            loss = criterion(zs, ts)
            loss.backward()
            optimizer.step()
            optimizer.zero grad()
            iters.append(n)
            losses.append(float(loss)/batch size)
            # if n%1 == 0:
            n += 1
       iters sub.append(epoch)
        train_cost = float(loss.detach().numpy())
        train_pos_acc, train_neg_acc = get_accuracy(model, t_data, batch_size)
       valid_pos_acc, valid_neg_acc = get_accuracy(model, v_data, batch_size)
        t_pos.append(train_pos_acc)
        t negs.append(train neg acc)
        v pos.append(valid pos acc)
       v negs.append(valid neg acc)
        print ("Iter %d. [Train Pos Acc %.0f%%, Train Neg Acc %.0f%%, Valid Pos Acc %.0f%%, Valid Ne
g Acc %.0f%%, Loss %f]" % (
              n, train pos acc * 100, train neg acc * 100, valid pos acc * 100, valid neg acc * 100,
train cost))
       now = time.time()
        print("last iteration: "+str(round(now-last, 5))+" seconds")
        print("elapsed time: "+str(round(now-start, 5))+" seconds\n")
        last = now
        if (checkpoint_path is not None):
            torch.save(model.state_dict(), checkpoint_path.format(n))
    return iters, losses, iters_sub, t_pos, t_negs, v_pos, v_negs
def plot learning curve (iters, losses, iters sub, t pos, t negs, v pos, v negs):
    plt.title("Learning Curve: Loss per Iteration")
    plt.plot(iters, losses, label="Train")
   plt.xlabel("Iteration")
   plt.ylabel("loss")
   plt.show()
    plt.title("Learning Curve: Accuracy per Epoch")
    plt.plot(iters sub, t pos, label="Training Positive")
    plt.plot(iters_sub, t_negs, label="Training Negative")
    plt.plot(iters_sub, v_pos, label="Validation Positive")
    plt.plot(iters sub, v negs, label="Validation Negative")
    plt.xlabel("Epoch")
    plt.ylabel("Accuracy")
    plt.legend(loc='best')
    plt.show()
                                                                                                 I
```

# Part (b) -- 2 pts

Sanity check your code from Q3(a) and from Q2(a) and Q2(b) by showing that your models can memorize a very small subset of the training set (e.g. 5 images). You should be able to achieve 90%+ accuracy relatively quickly (within ~30 or so iterations).

```
In [28]:
# Write your code here. Remember to include your results so that your TA can
# see that your model attains a high training accuracy. (UPDATED March 12)
learning curve info = train model(mod, train data[:2, :, :, :, :], valid_data, batch_size=6, num
epochs=60, checkpoint path='./overfit/ckpt-overfit-epoch{}.pk') # look! it works!
plot learning curve(*learning curve info)
Iter 1. [Train Pos Acc 100%, Train Neg Acc 0%, Valid Pos Acc 100%, Valid Neg Acc 0%, Loss 0.693051
last iteration: 2.13539 seconds
elapsed time: 2.13539 seconds
Iter 2. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.692960
last iteration: 2.17033 seconds
elapsed time: 4.30572 seconds
Iter 3. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.693251
last iteration: 2.13687 seconds
elapsed time: 6.44259 seconds
Iter 4. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.693574
last iteration: 2.1046 seconds
elapsed time: 8.54719 seconds
Iter 5. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.693917
last iteration: 2.18444 seconds
elapsed time: 10.73163 seconds
Iter 6. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.693470
last iteration: 2.30609 seconds
elapsed time: 13.03772 seconds
Iter 7. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.693524
last iteration: 2.40067 seconds
elapsed time: 15.43839 seconds
Iter 8. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.692632
last iteration: 2.24154 seconds
elapsed time: 17.67992 seconds
Iter 9. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.692293
last iteration: 2.28666 seconds
elapsed time: 19.96658 seconds
Iter 10. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.69249
last iteration: 2.23634 seconds
elapsed time: 22.20292 seconds
Iter 11. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.68646
last iteration: 2.44575 seconds
elapsed time: 24.64867 seconds
Iter 12. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.67217
last iteration: 2.49473 seconds
elapsed time: 27.1434 seconds
Iter 13. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.66246
61
```

last iteration: 2.43302 seconds elapsed time: 29.57641 seconds

```
Iter 14. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.65460
0]
last iteration: 2.6824 seconds
elapsed time: 32.25881 seconds
Iter 15. [Train Pos Acc 67%, Train Neg Acc 67%, Valid Pos Acc 3%, Valid Neg Acc 97%, Loss 0.701975
last iteration: 2.31876 seconds
elapsed time: 34.57757 seconds
Iter 16. [Train Pos Acc 67%, Train Neg Acc 67%, Valid Pos Acc 3%, Valid Neg Acc 97%, Loss 0.671467
last iteration: 2.24968 seconds
elapsed time: 36.82725 seconds
Iter 17. [Train Pos Acc 67%, Train Neg Acc 67%, Valid Pos Acc 11%, Valid Neg Acc 94%, Loss 0.55920
last iteration: 2.25463 seconds
elapsed time: 39.08189 seconds
Iter 18. [Train Pos Acc 83%, Train Neg Acc 83%, Valid Pos Acc 17%, Valid Neg Acc 94%, Loss 0.55772
last iteration: 2.56101 seconds
elapsed time: 41.6429 seconds
Iter 19. [Train Pos Acc 83%, Train Neg Acc 67%, Valid Pos Acc 39%, Valid Neg Acc 78%, Loss 0.57176
last iteration: 2.32105 seconds
elapsed time: 43.96395 seconds
Iter 20. [Train Pos Acc 67%, Train Neg Acc 67%, Valid Pos Acc 33%, Valid Neg Acc 81%, Loss 0.54229
last iteration: 2.52055 seconds
elapsed time: 46.4845 seconds
Iter 21. [Train Pos Acc 67%, Train Neg Acc 67%, Valid Pos Acc 42%, Valid Neg Acc 81%, Loss 0.59250
last iteration: 2.16862 seconds
elapsed time: 48.65312 seconds
Iter 22. [Train Pos Acc 83%, Train Neg Acc 83%, Valid Pos Acc 50%, Valid Neg Acc 75%, Loss 0.56197
last iteration: 2.28346 seconds
elapsed time: 50.93658 seconds
Iter 23. [Train Pos Acc 83%, Train Neg Acc 83%, Valid Pos Acc 67%, Valid Neg Acc 69%, Loss 0.58994
41
last iteration: 2.76705 seconds
elapsed time: 53.70363 seconds
Iter 24. [Train Pos Acc 83%, Train Neg Acc 83%, Valid Pos Acc 69%, Valid Neg Acc 69%, Loss 0.50717
21
last iteration: 8.14746 seconds
elapsed time: 61.85109 seconds
Iter 25. [Train Pos Acc 83%, Train Neg Acc 83%, Valid Pos Acc 67%, Valid Neg Acc 72%, Loss 0.39576
01
last iteration: 6.65583 seconds
elapsed time: 68.50692 seconds
Iter 26. [Train Pos Acc 83%, Train Neg Acc 83%, Valid Pos Acc 64%, Valid Neg Acc 67%, Loss 0.32469
last iteration: 5.37052 seconds
elapsed time: 73.87744 seconds
Iter 27. [Train Pos Acc 100%, Train Neg Acc 83%, Valid Pos Acc 78%, Valid Neg Acc 61%, Loss 0.4155
last iteration: 3.59599 seconds
elapsed time: 77.47343 seconds
Iter 28. [Train Pos Acc 100%, Train Neg Acc 67%, Valid Pos Acc 94%, Valid Neg Acc 42%, Loss 0.2185
last iteration: 2.34622 seconds
elapsed time: 79.81965 seconds
```

Iter 29. [Train Pos Acc 100%, Train Neg Acc 50%, Valid Pos Acc 100%, Valid Neg Acc 22%, Loss 0.135

```
9411
```

last iteration: 2.23963 seconds elapsed time: 82.05929 seconds

Iter 30. [Train Pos Acc 100%, Train Neg Acc 50%, Valid Pos Acc 100%, Valid Neg Acc 31%, Loss 0.321

last iteration: 2.20212 seconds elapsed time: 84.26141 seconds

Iter 31. [Train Pos Acc 100%, Train Neg Acc 83%, Valid Pos Acc 92%, Valid Neg Acc 47%, Loss 0.1895

pT]

last iteration: 2.31766 seconds elapsed time: 86.57907 seconds

Iter 32. [Train Pos Acc 50%, Train Neg Acc 100%, Valid Pos Acc 69%, Valid Neg Acc 61%, Loss 0.3297

last iteration: 2.86142 seconds elapsed time: 89.44049 seconds

Iter 33. [Train Pos Acc 67%, Train Neg Acc 100%, Valid Pos Acc 69%, Valid Neg Acc 61%, Loss 0.2505

last iteration: 3.13182 seconds elapsed time: 92.57231 seconds

Iter 34. [Train Pos Acc 100%, Train Neg Acc 83%, Valid Pos Acc 86%, Valid Neg Acc 61%, Loss 0.1753

last iteration: 2.10147 seconds elapsed time: 94.67378 seconds

Iter 35. [Train Pos Acc 100%, Train Neg Acc 83%, Valid Pos Acc 89%, Valid Neg Acc 53%, Loss 0.2192 49]

last iteration: 2.22267 seconds elapsed time: 96.89646 seconds

Iter 36. [Train Pos Acc 100%, Train Neg Acc 83%, Valid Pos Acc 89%, Valid Neg Acc 50%, Loss 0.2843

last iteration: 2.41731 seconds elapsed time: 99.31377 seconds

Iter 37. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 86%, Valid Neg Acc 58%, Loss 0.323 604]

last iteration: 2.64689 seconds elapsed time: 101.96066 seconds

Iter 38. [Train Pos Acc 83%, Train Neg Acc 100%, Valid Pos Acc 83%, Valid Neg Acc 64%, Loss 0.0571

last iteration: 2.37409 seconds elapsed time: 104.33474 seconds

Iter 39. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 89%, Valid Neg Acc 58%, Loss 0.230 356]

last iteration: 2.37454 seconds elapsed time: 106.70929 seconds

Iter 40. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 89%, Valid Neg Acc 47%, Loss 0.056

last iteration: 2.20879 seconds elapsed time: 108.91807 seconds

Iter 41. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 39%, Loss 0.135

last iteration: 2.2217 seconds elapsed time: 111.13977 seconds

Iter 42. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 36%, Loss 0.091
241]

last iteration: 2.36227 seconds elapsed time: 113.50204 seconds

Iter 43. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 25%, Loss 0.036 236]

last iteration: 2.09435 seconds elapsed time: 115.59639 seconds

Iter 44. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 25%, Loss 0.011

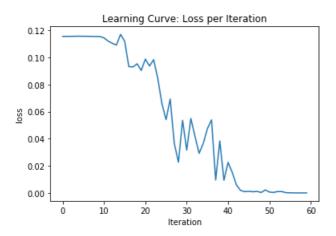
last iteration: 2.09605 seconds

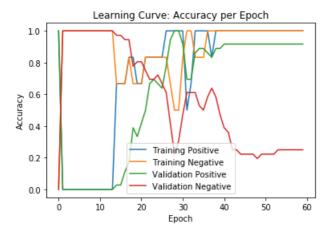
elapsed time: 117.69244 seconds Iter 45. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 22%, Loss 0.005 last iteration: 2.08954 seconds elapsed time: 119.78198 seconds Iter 46. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 22%, Loss 0.007 last iteration: 2.08049 seconds elapsed time: 121.86247 seconds Iter 47. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 22%, Loss 0.005 last iteration: 2.26244 seconds elapsed time: 124.12491 seconds Iter 48. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 22%, Loss 0.006 last iteration: 2.20561 seconds elapsed time: 126.33052 seconds Iter 49. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 19%, Loss 0.001 last iteration: 2.07839 seconds elapsed time: 128.40891 seconds Iter 50. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 22%, Loss 0.013 last iteration: 2.00285 seconds elapsed time: 130.41176 seconds Iter 51. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 22%, Loss 0.003 7651 last iteration: 2.03822 seconds elapsed time: 132.44998 seconds Iter 52. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 22%, Loss 0.002 last iteration: 2.0847 seconds elapsed time: 134.53467 seconds Iter 53. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 22%, Loss 0.006 last iteration: 2.04741 seconds elapsed time: 136.58209 seconds Iter 54. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 25%, Loss 0.006 last iteration: 2.0697 seconds elapsed time: 138.65179 seconds Iter 55. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 25%, Loss 0.001 0011 last iteration: 2.28229 seconds elapsed time: 140.93407 seconds Iter 56. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 25%, Loss 0.000 7241 last iteration: 2.59143 seconds elapsed time: 143.52551 seconds Iter 57. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 25%, Loss 0.000 0571 last iteration: 2.33089 seconds elapsed time: 145.8564 seconds Iter 58. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 25%, Loss 0.000 last iteration: 2.11706 seconds elapsed time: 147.97346 seconds Iter 59. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 25%, Loss 0.000

last iteration: 2.14607 seconds elapsed time: 150.11954 seconds

Iter 60. [Train Pos Acc 100%, Train Neg Acc 100%, Valid Pos Acc 92%, Valid Neg Acc 25%, Loss 0.000 024]

last iteration: 2.08795 seconds elapsed time: 152.20748 seconds





# Part (c) -- 4 pts

Train your models from Q2(a) and Q2(b). You will want to explore the effects of a few hyperparameters, including the learning rate, batch size, choice of n, and potentially the kernel size. You do not need to check all values for all hyperparameters. Instead, get an intuition about what each of the parameters do.

In this section, explain how you tuned your hyperparameters.

### In [29]:

```
# Include the training curves for the two models.
mod1 = CNNChannel(6)
learning curve info1 = train model(mod1, train data, valid data, batch size=50, num epochs=15, chec
kpoint path='./channel/ckpt1-epoch{}.pk')
mod2 = CNNChannel(8)
learning_curve_info2 = train_model(mod2, train_data, valid_data, batch_size=30, num_epochs=10, chec
kpoint_path='./channel/ckpt2-epoch{}.pk')
mod3 = CNNChannel()
learning curve info3 = train model(mod3, train data, valid data, batch size=30, learning rate=0.000
1, num epochs=10, checkpoint path='./channel/ckpt3-epoch{}.pk')
mod4 = CNN(6)
learning curve info4 = train model(mod4, train data, valid data, batch size=50, num epochs=15, chec
kpoint_path='./vanilla/ckpt4-epoch{}.pk')
mod5 = CNN(8)
learning_curve_info5 = train_model(mod5, train_data, valid_data, batch_size=30, num_epochs=10, chec
kpoint path='./vanilla/ckpt5-epoch{}.pk')
mod6 = CNN()
learning curve info6 = train model (mode train data valid data batch circe=30 learning rate=0 000
```

```
realisting curve intoo - crain moder(modo, crain data, varid data, bacch Size-30, realisting rate-0.000
1, num_epochs=10, checkpoint_path='./vanilla/ckpt6-epoch{}.pk')
11 11 11
In this case we simply tried small variations of single paramaters to see what would happen and de
which paramaters would lead to a more accurate model. While our low learning rate model was the be
took much longer to train. We know increasing the number of epochs will generally increase the tra
accuracy at the cost of a potential overfit, so we tried varying the number of epochs less than ot
parameters. We also found that increasing the value of the hyperparameter n makes the model perfor
m better.
Note: the number of epochs we used was fairly small; even for the 15 epoch case we could have trai
the model more and achieved higher accuracies which is implied by our training curves for models
mod3 and mod4 which were the two models with lower learning rates.
Iter 6. [Train Pos Acc 84%, Train Neg Acc 27%, Valid Pos Acc 86%, Valid Neg Acc 25%, Loss 0.694546
last iteration: 72.11294 seconds
elapsed time: 72.11294 seconds
Iter 12. [Train Pos Acc 98%, Train Neg Acc 24%, Valid Pos Acc 100%, Valid Neg Acc 31%, Loss 0.6692
681
last iteration: 64.7012 seconds
elapsed time: 136.81414 seconds
Iter 18. [Train Pos Acc 75%, Train Neg Acc 71%, Valid Pos Acc 86%, Valid Neg Acc 78%, Loss 0.41447
last iteration: 68.83449 seconds
elapsed time: 205.64863 seconds
Iter 24. [Train Pos Acc 93%, Train Neg Acc 56%, Valid Pos Acc 97%, Valid Neg Acc 61%, Loss 0.44722
last iteration: 72.87511 seconds
elapsed time: 278.52374 seconds
Iter 30. [Train Pos Acc 91%, Train Neg Acc 68%, Valid Pos Acc 97%, Valid Neg Acc 69%, Loss 0.31885
last iteration: 78.22578 seconds
elapsed time: 356.74952 seconds
Iter 36. [Train Pos Acc 91%, Train Neg Acc 72%, Valid Pos Acc 97%, Valid Neg Acc 69%, Loss 0.52454
81
last iteration: 58.33639 seconds
elapsed time: 415.08591 seconds
Iter 42. [Train Pos Acc 93%, Train Neg Acc 72%, Valid Pos Acc 94%, Valid Neg Acc 69%, Loss 0.41067
61
last iteration: 58.01516 seconds
elapsed time: 473.10107 seconds
Iter 48. [Train Pos Acc 94%, Train Neg Acc 72%, Valid Pos Acc 94%, Valid Neg Acc 67%, Loss 0.32186
last iteration: 56.64903 seconds
elapsed time: 529.75011 seconds
Iter 54. [Train Pos Acc 89%, Train Neg Acc 82%, Valid Pos Acc 92%, Valid Neg Acc 72%, Loss 0.40569
last iteration: 55.94641 seconds
elapsed time: 585.69652 seconds
Iter 60. [Train Pos Acc 75%, Train Neg Acc 88%, Valid Pos Acc 86%, Valid Neg Acc 83%, Loss 0.35065
51
last iteration: 57.71069 seconds
elapsed time: 643.40721 seconds
Iter 66. [Train Pos Acc 94%, Train Neg Acc 67%, Valid Pos Acc 97%, Valid Neg Acc 64%, Loss 0.44883
61
last iteration: 63.35966 seconds
elapsed time: 706.76687 seconds
Iter 72. [Train Pos Acc 92%, Train Neg Acc 73%, Valid Pos Acc 92%, Valid Neg Acc 69%, Loss 0.45855
```

```
1]
last iteration: 72.83219 seconds
elapsed time: 779.59906 seconds
Iter 78. [Train Pos Acc 84%, Train Neg Acc 86%, Valid Pos Acc 89%, Valid Neg Acc 75%, Loss 0.53234
last iteration: 58.0984 seconds
elapsed time: 837.69746 seconds
Iter 84. [Train Pos Acc 93%, Train Neg Acc 78%, Valid Pos Acc 94%, Valid Neg Acc 75%, Loss 0.34731
last iteration: 56.49434 seconds
elapsed time: 894.1918 seconds
Iter 90. [Train Pos Acc 91%, Train Neg Acc 83%, Valid Pos Acc 92%, Valid Neg Acc 75%, Loss 0.38935
21
last iteration: 55.95836 seconds
elapsed time: 950.15017 seconds
Iter 10. [Train Pos Acc 72%, Train Neg Acc 34%, Valid Pos Acc 86%, Valid Neg Acc 22%, Loss 0.69281
last iteration: 61.19633 seconds
elapsed time: 61.19633 seconds
Iter 20. [Train Pos Acc 31%, Train Neg Acc 85%, Valid Pos Acc 22%, Valid Neg Acc 94%, Loss 0.68767
last iteration: 58.64419 seconds
elapsed time: 119.84053 seconds
Iter 30. [Train Pos Acc 77%, Train Neg Acc 73%, Valid Pos Acc 86%, Valid Neg Acc 78%, Loss 0.54230
last iteration: 58.28416 seconds
elapsed time: 178.12469 seconds
Iter 40. [Train Pos Acc 94%, Train Neg Acc 65%, Valid Pos Acc 97%, Valid Neg Acc 61%, Loss 0.51704
last iteration: 58.77492 seconds
elapsed time: 236.89961 seconds
Iter 50. [Train Pos Acc 96%, Train Neg Acc 65%, Valid Pos Acc 94%, Valid Neg Acc 64%, Loss 0.47235
last iteration: 58.34205 seconds
elapsed time: 295.24165 seconds
Iter 60. [Train Pos Acc 86%, Train Neg Acc 81%, Valid Pos Acc 92%, Valid Neg Acc 81%, Loss 0.55575
last iteration: 59.74434 seconds
elapsed time: 354.986 seconds
Iter 70. [Train Pos Acc 95%, Train Neg Acc 71%, Valid Pos Acc 94%, Valid Neg Acc 67%, Loss 0.34457
last iteration: 69.79438 seconds
elapsed time: 424.78038 seconds
Iter 80. [Train Pos Acc 79%, Train Neg Acc 90%, Valid Pos Acc 83%, Valid Neg Acc 86%, Loss 0.30453
81
last iteration: 71.88166 seconds
elapsed time: 496.66204 seconds
Iter 90. [Train Pos Acc 88%, Train Neg Acc 83%, Valid Pos Acc 94%, Valid Neg Acc 78%, Loss 0.46890
6]
last iteration: 68.2665 seconds
elapsed time: 564.92853 seconds
Iter 100. [Train Pos Acc 96%, Train Neg Acc 67%, Valid Pos Acc 97%, Valid Neg Acc 67%, Loss 0.2493
last iteration: 66.36134 seconds
elapsed time: 631.28987 seconds
Iter 10. [Train Pos Acc 100%, Train Neg Acc 0%, Valid Pos Acc 100%, Valid Neg Acc 0%, Loss 0.69330
81
last iteration: 61.90752 seconds
elapsed time: 61.90752 seconds
```

Iter 20. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.69309

```
9]
last iteration: 66.10841 seconds
elapsed time: 128.01593 seconds
Iter 30. [Train Pos Acc 90%, Train Neg Acc 29%, Valid Pos Acc 81%, Valid Neg Acc 36%, Loss 0.69309
last iteration: 61.64907 seconds
elapsed time: 189.665 seconds
Iter 40. [Train Pos Acc 42%, Train Neg Acc 81%, Valid Pos Acc 25%, Valid Neg Acc 83%, Loss 0.69232
last iteration: 65.61997 seconds
elapsed time: 255.28498 seconds
Iter 50. [Train Pos Acc 52%, Train Neg Acc 79%, Valid Pos Acc 33%, Valid Neg Acc 81%, Loss 0.69258
last iteration: 67.37093 seconds
elapsed time: 322.65591 seconds
Iter 60. [Train Pos Acc 86%, Train Neg Acc 66%, Valid Pos Acc 81%, Valid Neg Acc 61%, Loss 0.68764
last iteration: 64.91101 seconds
elapsed time: 387.56692 seconds
Iter 70. [Train Pos Acc 83%, Train Neg Acc 65%, Valid Pos Acc 83%, Valid Neg Acc 64%, Loss 0.67112
last iteration: 62.92577 seconds
elapsed time: 450.49269 seconds
Iter 80. [Train Pos Acc 69%, Train Neg Acc 79%, Valid Pos Acc 61%, Valid Neg Acc 78%, Loss 0.61561
last iteration: 57.40863 seconds
elapsed time: 507.90132 seconds
Iter 90. [Train Pos Acc 80%, Train Neg Acc 77%, Valid Pos Acc 89%, Valid Neg Acc 75%, Loss 0.54600
11
last iteration: 57.70981 seconds
elapsed time: 565.61113 seconds
Iter 100. [Train Pos Acc 67%, Train Neg Acc 89%, Valid Pos Acc 75%, Valid Neg Acc 92%, Loss 0.5821
last iteration: 57.22907 seconds
elapsed time: 622.8402 seconds
Iter 6. [Train Pos Acc 100%, Train Neg Acc 0%, Valid Pos Acc 100%, Valid Neg Acc 0%, Loss 0.695081
last iteration: 60.88317 seconds
elapsed time: 60.88317 seconds
Iter 12. [Train Pos Acc 3%, Train Neg Acc 97%, Valid Pos Acc 3%, Valid Neg Acc 100%, Loss 0.692903
last iteration: 67.52001 seconds
elapsed time: 128.40319 seconds
Iter 18. [Train Pos Acc 100%, Train Neg Acc 0%, Valid Pos Acc 100%, Valid Neg Acc 0%, Loss 0.69324
11
last iteration: 71.73423 seconds
elapsed time: 200.13742 seconds
Iter 24. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.69327
11
last iteration: 72.7865 seconds
elapsed time: 272.92391 seconds
Iter 30. [Train Pos Acc 59%, Train Neg Acc 39%, Valid Pos Acc 53%, Valid Neg Acc 56%, Loss 0.69316
last iteration: 76.06064 seconds
elapsed time: 348.98455 seconds
Iter 36. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.69313
last iteration: 75.31463 seconds
elapsed time: 424.29918 seconds
Iter 42. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.69310
81
```

last iteration: 71.6824 seconds elapsed time: 495.98158 seconds Iter 48. [Train Pos Acc 45%, Train Neg Acc 55%, Valid Pos Acc 25%, Valid Neg Acc 75%, Loss 0.69317 last iteration: 70.84885 seconds elapsed time: 566.83043 seconds Iter 54. [Train Pos Acc 48%, Train Neg Acc 55%, Valid Pos Acc 25%, Valid Neg Acc 69%, Loss 0.69315 last iteration: 71.80815 seconds elapsed time: 638.63858 seconds Iter 60. [Train Pos Acc 80%, Train Neg Acc 18%, Valid Pos Acc 83%, Valid Neg Acc 19%, Loss 0.69326 last iteration: 68.94164 seconds elapsed time: 707.58022 seconds Iter 66. [Train Pos Acc 82%, Train Neg Acc 25%, Valid Pos Acc 86%, Valid Neg Acc 22%, Loss 0.69399 81 last iteration: 62.97604 seconds elapsed time: 770.55625 seconds Iter 72. [Train Pos Acc 32%, Train Neg Acc 81%, Valid Pos Acc 28%, Valid Neg Acc 92%, Loss 0.69123 last iteration: 60.99845 seconds elapsed time: 831.5547 seconds Iter 78. [Train Pos Acc 81%, Train Neg Acc 54%, Valid Pos Acc 83%, Valid Neg Acc 53%, Loss 0.69971 last iteration: 60.48036 seconds elapsed time: 892.03506 seconds Iter 84. [Train Pos Acc 70%, Train Neg Acc 67%, Valid Pos Acc 75%, Valid Neg Acc 53%, Loss 0.67007 last iteration: 59.847 seconds elapsed time: 951.88206 seconds Iter 90. [Train Pos Acc 59%, Train Neg Acc 78%, Valid Pos Acc 72%, Valid Neg Acc 72%, Loss 0.56603 last iteration: 59.6022 seconds elapsed time: 1011.48426 seconds Iter 10. [Train Pos Acc 84%, Train Neg Acc 15%, Valid Pos Acc 94%, Valid Neg Acc 6%, Loss 0.710438 last iteration: 64.44274 seconds elapsed time: 64.44274 seconds Iter 20. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.69243 last iteration: 63.51025 seconds elapsed time: 127.95299 seconds Iter 30. [Train Pos Acc 5%, Train Neg Acc 97%, Valid Pos Acc 3%, Valid Neg Acc 97%, Loss 0.691852] last iteration: 64.03781 seconds elapsed time: 191.9908 seconds Iter 40. [Train Pos Acc 51%, Train Neg Acc 51%, Valid Pos Acc 50%, Valid Neg Acc 50%, Loss 0.68922 last iteration: 64.11956 seconds elapsed time: 256.11035 seconds Iter 50. [Train Pos Acc 100%, Train Neg Acc 1%, Valid Pos Acc 100%, Valid Neg Acc 0%, Loss 0.69193 last iteration: 63.22794 seconds elapsed time: 319.33829 seconds Iter 60. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.69207 last iteration: 64.52148 seconds elapsed time: 383.85977 seconds Iter 70. [Train Pos Acc 98%, Train Neg Acc 3%, Valid Pos Acc 97%, Valid Neg Acc 3%, Loss 0.694078] last iteration: 62.67442 seconds elapsed time: 446.53419 seconds

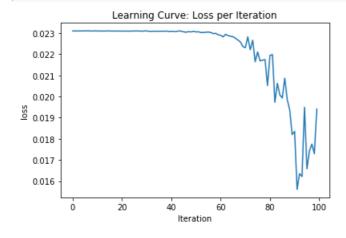
```
Iter 80. [Train Pos Acc 14%, Train Neg Acc 88%, Valid Pos Acc 28%, Valid Neg Acc 75%, Loss 0.69380
last iteration: 63.12421 seconds
elapsed time: 509.65841 seconds
Iter 90. [Train Pos Acc 14%, Train Neg Acc 87%, Valid Pos Acc 11%, Valid Neg Acc 92%, Loss 0.69222
last iteration: 62.90979 seconds
elapsed time: 572.5682 seconds
Iter 100. [Train Pos Acc 3%, Train Neg Acc 95%, Valid Pos Acc 11%, Valid Neg Acc 89%, Loss 0.69311
last iteration: 62.53878 seconds
elapsed time: 635.10698 seconds
Iter 10. [Train Pos Acc 0%, Train Neg Acc 100%, Valid Pos Acc 0%, Valid Neg Acc 100%, Loss 0.69463
last iteration: 58.33911 seconds
elapsed time: 58.33911 seconds
Iter 20. [Train Pos Acc 100%, Train Neg Acc 0%, Valid Pos Acc 100%, Valid Neg Acc 0%, Loss 0.69368
last iteration: 60.14797 seconds
elapsed time: 118.48709 seconds
Iter 30. [Train Pos Acc 3%, Train Neg Acc 99%, Valid Pos Acc 8%, Valid Neg Acc 97%, Loss 0.693185]
last iteration: 58.34899 seconds
elapsed time: 176.83607 seconds
Iter 40. [Train Pos Acc 59%, Train Neg Acc 40%, Valid Pos Acc 78%, Valid Neg Acc 17%, Loss 0.69284
last iteration: 61.25033 seconds
elapsed time: 238.0864 seconds
Iter 50. [Train Pos Acc 20%, Train Neg Acc 83%, Valid Pos Acc 19%, Valid Neg Acc 83%, Loss 0.69378
last iteration: 57.94207 seconds
elapsed time: 296.02847 seconds
Iter 60. [Train Pos Acc 43%, Train Neg Acc 57%, Valid Pos Acc 47%, Valid Neg Acc 53%, Loss 0.69283
last iteration: 58.78681 seconds
elapsed time: 354.81529 seconds
Iter 70. [Train Pos Acc 46%, Train Neg Acc 54%, Valid Pos Acc 53%, Valid Neg Acc 47%, Loss 0.69349
last iteration: 59.30454 seconds
elapsed time: 414.11983 seconds
Iter 80. [Train Pos Acc 90%, Train Neg Acc 8%, Valid Pos Acc 92%, Valid Neg Acc 8%, Loss 0.694216]
last iteration: 59.41922 seconds
elapsed time: 473.53905 seconds
Iter 90. [Train Pos Acc 12%, Train Neg Acc 88%, Valid Pos Acc 8%, Valid Neg Acc 94%, Loss 0.694192
last iteration: 58.75799 seconds
elapsed time: 532.29704 seconds
Iter 100. [Train Pos Acc 88%, Train Neg Acc 11%, Valid Pos Acc 92%, Valid Neg Acc 8%, Loss 0.69393
last iteration: 58.95736 seconds
elapsed time: 591.25439 seconds
```

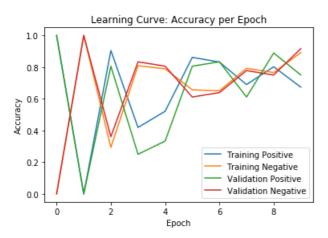
## Part (d) -- 2 pts


Include your training curves for the **best** models from each of Q2(a) and Q2(b). These are the models that you will use in Question 4.

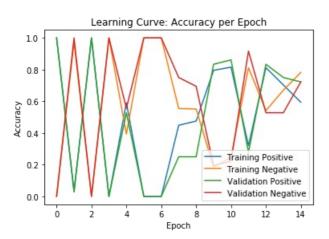
In [32]:

```
# Include the training curves for the two models.
plot_learning_curve(*learning_curve_info3)
plot_learning_curve(*learning_curve_info4)
```









# Question 4.

## Part (a) -- 3 pts

Report the test accuracies of your **single best** model, separately for the two test sets. Do this by choosing the checkpoint of the model architecture that produces the best validation accuracy. That is, if your model attained the best validation accuracy in epoch 12, then the weights at epoch 12 is what you should be using to report the test accuracy.

#### In [60]:

[Test Pos Acc M 67%, Test Neg Acc M 90%, Test Pos Acc W 67%, Test Neg Acc W 93%]

#### Out[60]:

'\nLoading a checkpoint was not required for our model because it had achieved the highest validat ion accuracy\nupon the last epoch. $\n'$ 

# Part (b) -- 2 pts

Display one set of men's shoes that your model correctly classified as being from the same pair.

If your test accuracy was not 100% on the men's shoes test set, display one set of inputs that your model classified incorrectly.

## In [51]:

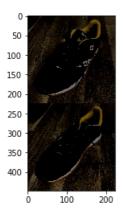
```
same, dif = make_data(test_m)
pred = model(torch.Tensor(same[0]).unsqueeze(0).transpose(1, 3)).max(1, keepdim=True)[1]
print(bool(pred[0])) # poggers
plt.imshow(same[0])

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).
```

True

## Out[51]:

<matplotlib.image.AxesImage at 0x2788460c748>



# Part (c) -- 2 pts

Display one set of women's shoes that your model correctly classified as being from the same pair

Display one set of noments shows that your model contoury diagonica actioning from the came pain.

If your test accuracy was not 100% on the women's shoes test set, display one set of inputs that your model classified incorrectly.

### In [52]:

```
same, dif = make_data(test_w)
pred = model(torch.Tensor(same[0]).unsqueeze(0).transpose(1, 3)).max(1, keepdim=True)[1]
print(bool(pred[0])) # poggers
plt.imshow(same[0])
Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for
```

True

## Out[52]:

integers).

<matplotlib.image.AxesImage at 0x278847984a8>

